

What the constraint leaves behind
early-layer entropy stays low at 5× longer training

